

On The Expressiveness Of Masked Hard-Attention Transformers

Course in Foundations of Neural Networks

BORELLI ROBERTO
`borelli.roberto@spes.uniud.it`

University of Udine

January 30, 2025

Abstract

This work presents the paper by Yang et al. which characterizes the expressiveness of a particular class of transformers with hard attention, where attention is focused on exactly one position at a time. It is shown how to compile a transformer model into the language B-RASP. Furthermore, it is established that B-RASP is equivalent to star-free languages. The proof proceeds in two directions, employing two distinct characterizations of star-free languages: linear temporal logic over finite traces and cascades of reset automata. This study offers a deeper understanding of the transformer formalism, revealing that (i) both the feed-forward and self-attention sublayers play crucial roles, and (ii) increasing the number of layers in a transformer enhances its expressive power. This latter result contrasts with the universal approximation theorem for standard feedforward neural networks.

Contents

1 Introduction

- Transformers
- B-RASP

2 Transformers and B-RASP

- Transformers \rightarrow B-RASP
- Transformers \leftarrow B-RASP

3 B-RASP and Star Free Languages

- B-RASP \rightarrow LTL
- B-RASP \leftarrow Automata Cascades

4 Conclusions

Introduction

Transformers

B-RASP

Transformer: assumptions


The Transformer studied here:

- is an encoder-only model
- uses **unique hard attention**: all attention is focused on exactly one position

Further simplifications:


- only single-head attention
- no layer normalization
- no positional embeddings

Transformer: general view




- $T(w)$ is the output on input w
- $[T(w)]_i \in \mathbb{R}^d$ is the i -th output vector

Transformer: general view




- $T(w)$ is the output on input w
- $[T(w)]_i \in \mathbb{R}^d$ is the i -th output vector

Transformer: general view




- $T(w)$ is the output on input w
- $[T(w)]_i \in \mathbb{R}^d$ is the i -th output vector

Transformer: general view



- $T(w)$ is the output on input w
- $[T(w)]_i \in \mathbb{R}^d$ is the i -th output vector
- A word w is accepted iff the linear projection of $[T(w)]_n$ is nonnegative


Transformer: layer



$$Layer_j : (\mathbb{R}^d)^+ \rightarrow (\mathbb{R}^d)^+$$

$$(x_1, \dots, x_n) \rightarrow (y_1, \dots, y_n)$$

Transformer: layer




FeedForward SubLayer

It's a feed forward network with two layers and ReLU activation

$$FF(x_i) = (\max \{x_i \cdot W_1 + b_1, 0\}) W_2 + b_2$$

Transformer: layer



Self Attention SubLayer

- score function: bilinear function $f_S : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$
- value function: linear transformation $f_V : \mathbb{R}^d \rightarrow \mathbb{R}^d$
- mask: one of $M(i, j) = 1$; $M(i, j) = i < j$; $M(i, j) = j < i$
- selector function: either $f_C = \max$ or $f_C = \min$


Example of future-masked leftmost-hard attention

Let $M(i, j) = j < i$ and $f_C = \min$. Let x_1, \dots, x_5 be input vectors. The steps for $SA(x_4)$ are the following.




Example of future-masked leftmost-hard attention

Let $M(i, j) = j < i$ and $f_C = \min$. Let x_1, \dots, x_5 be input vectors. The steps for $SA(x_4)$ are the following.



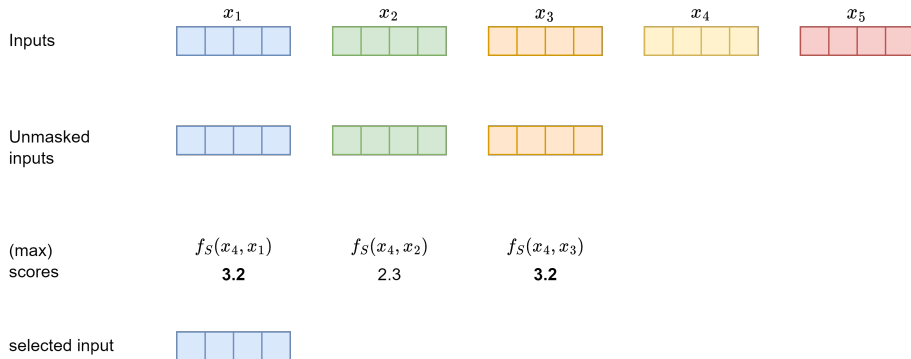
Example of future-masked leftmost-hard attention

Let $M(i, j) = j < i$ and $f_C = \min$. Let x_1, \dots, x_5 be input vectors. The steps for $SA(x_4)$ are the following.




Example of future-masked leftmost-hard attention

Let $M(i, j) = j < i$ and $f_C = \min$. Let x_1, \dots, x_5 be input vectors. The steps for $SA(x_4)$ are the following.



Example of future-masked leftmost-hard attention

Let $M(i, j) = j < i$ and $f_C = \min$. Let x_1, \dots, x_5 be input vectors. The steps for $SA(x_4)$ are the following.



B-RASP: definition

- It's a programming language intended to help *think like transformers*

B-RASP: definition

- It's a programming language intended to help *think like transformers*
- A B-RASP program p is a sequence of definitions of boolean vectors
- Let $\Sigma = \{\sigma_1, \dots, \sigma_m\}$ be an alphabet and $w = w_1 \dots w_n$ a word in Σ^*

B-RASP: definition

- It's a programming language intended to help *think like transformers*
- A B-RASP program p is a sequence of definitions of boolean vectors
- Let $\Sigma = \{\sigma_1, \dots, \sigma_m\}$ be an alphabet and $w = w_1 \dots w_n$ a word in Σ^*
- **Initial vectors:** one-hot encoding of the input

B-RASP: definition

- It's a programming language intended to help *think like transformers*
- A B-RASP program p is a sequence of definitions of boolean vectors
- Let $\Sigma = \{\sigma_1, \dots, \sigma_m\}$ be an alphabet and $w = w_1 \dots w_n$ a word in Σ^*
- **Initial vectors:** $P_j(i) = 1 \leftrightarrow w_i = \sigma_j \quad j = 1, \dots, m$

B-RASP: definition

- It's a programming language intended to help *think like transformers*
- A B-RASP program p is a sequence of definitions of boolean vectors
- Let $\Sigma = \{\sigma_1, \dots, \sigma_m\}$ be an alphabet and $w = w_1 \dots w_n$ a word in Σ^*
- **Initial vectors:** $P_j(i) = 1 \leftrightarrow w_i = \sigma_j \quad j = 1, \dots, m$
- **New vectors:** for $k \geq m$ the vector P_{k+1} is defined combining the values of P_1, \dots, P_k with some operators
- **Output vector:** The last vector is the output vector denoted with Y

B-RASP: definition

- It's a programming language intended to help *think like transformers*
- A B-RASP program p is a sequence of definitions of boolean vectors
- Let $\Sigma = \{\sigma_1, \dots, \sigma_m\}$ be an alphabet and $w = w_1 \dots w_n$ a word in Σ^*
- **Initial vectors:** $P_j(i) = 1 \leftrightarrow w_i = \sigma_j \quad j = 1, \dots, m$
- **New vectors:** for $k \geq m$ the vector P_{k+1} is defined combining the values of P_1, \dots, P_k with some operators
- **Output vector:** The last vector is the output vector denoted with Y
- $w \in \mathcal{L}(p) \leftrightarrow Y(n) = 1$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

Syntax $P_{k+1}(i) := \text{selector}_j [\text{test}] V(i, j) : D(i)$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

Syntax $P_{k+1}(i) := \text{selector}_j [\text{test}] V(i, j) : D(i)$

Semantics

```

j' = select j such that test is satisfied
if j' == NULL
    Pk+1(i) = D(i)
else
    Pk+1(i) = V(i, j')
```

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

Syntax $P_{k+1}(i) := \text{min/max}_j [M(i,j) \wedge S(i,j)] V(i,j) : D(i)$

Semantics

```

j' = select min/max j such that  $M(i,j) \wedge S(i,j) = 1$ 
if j' == NULL
     $P_{k+1}(i) = D(i)$ 
else
     $P_{k+1}(i) = V(i, j')$ 

```

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : D(i)$$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

$$P_{k+1}(i) := \min/\max_j [\boxed{M(i,j)} \wedge S(i,j)] \quad V(i,j) : D(i)$$

Mask predicate

$$M(i,j) = \begin{cases} 1 & \text{no masking} \\ j < i & \text{future masking} \\ j > i & \text{past masking} \end{cases}$$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge \boxed{S(i,j)}] V(i,j) : D(i)$$

Score predicate

$$S(i,j) = \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\} \cup \{P_1(j), \dots, P_k(j)\})$$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \boxed{V(i,j)} : D(i)$$

Value predicate

$$V(i,j) = \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\} \cup \{P_1(j), \dots, P_k(j)\})$$

B-RASP: operations

New vectors can be defined in the following way.

Position-wise operations

$$P_{k+1}(i) := \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

Attention operations

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : \boxed{D(i)}$$

Default Value predicate

$$D(i) = \text{BoolCombination}(\{P_1(i), \dots, P_k(i)\})$$

B-RASP: example

Consider $\Sigma = \{a, b\}$, let Q_a and Q_b be the two input vectors. We define the following program p .

$$P_a(i) := \max_j [j < i, 1] Q_a(j) : 0$$

$$S_b(i) := \min_j [j > i, 1] Q_b(j) : 0$$

$$V := (Q_a \wedge S_b) \vee (Q_b \wedge P_a)$$

$$Y := \max_j [1, \neg V(j)] 0 : 1$$

It is easy to check that $\mathcal{L}(p) = (ab)^*$

B-RASP: normal form

Lemma (Normalization)

Any B-RASP program is equivalent to one in which all value predicates $V(i,j)$ and all score predicates $S(i,j)$ depend only on j .

Notice that the normalization procedure may require an exponential blowup.

Transformers and B-RASP

Transformers \rightarrow B-RASP

Transformers \leftarrow B-RASP

Transformers and computed values

Lemma

Let T be a transformer. Let \mathbb{F} be the union over all layers of the possible attention scores and components of the possible activation vectors. It holds that \mathbb{F} is finite.

Transformers and computed values

Lemma

Let T be a transformer. Let \mathbb{F} be the union over all layers of the possible attention scores and components of the possible activation vectors. It holds that \mathbb{F} is finite.

- Any attention score or component of an activation vector, can be represented with $B = \lceil \log_2 \mathbb{F} \rceil$ bits
- For a value v , let $\langle v \rangle_b$ be the bit corresponding to 2^b

Transformers and computed values

Lemma

Let T be a transformer. Let \mathbb{F} be the union over all layers of the possible attention scores and components of the possible activation vectors. It holds that \mathbb{F} is finite.

- Any attention score or component of an activation vector, can be represented with $B = \lceil \log_2 \mathbb{F} \rceil$ bits
- For a value v , let $\langle v \rangle_b$ be the bit corresponding to 2^b

Definition

A B-RASP program p simulates a transformer $T : \Sigma^+ \rightarrow (\mathbb{R}^d)^+$ if

- for each $b \in [B]$ and $k \in [d]$, p contains boolean vectors $Y_{k,b}$
- for every word w of length n , for $i \in [n]$, $b \in [B]$, $k \in [d]$, it holds $Y_{k,b}(i) = 1$ iff $\langle [T(w)]_{i,k} \rangle_b = 1$

Translating transformers to B-RASP

Lemma

For any masked hard-attention transformer T , there is a B-RASP program p_T that simulates T .

Translating transformers to B-RASP


Lemma

For any masked hard-attention transformer T , there is a B-RASP program p_T that simulates T .

Proof sketch

- The embedding function is simulated by

$$E_{k,b}(i) := \bigwedge_{\sigma \in \Sigma} Q_{\sigma}(i) \rightarrow \langle [Emb(\sigma)]_k \rangle_b$$



3-RASP program

The embedding function is simulated by

$$E_{k,b}(i) := \bigwedge_{\sigma \in \Sigma} Q_{\sigma}(i) \rightarrow \langle [Emb(\sigma)]_k \rangle_b$$

Translating transformers to B-RASP

Lemma

For any masked hard-attention transformer T , there is a B-RASP program p_T that simulates T .

Proof sketch

- The embedding function is simulated by
$$E_{k,b}(i) := \bigwedge_{\sigma \in \Sigma} Q_{\sigma}(i) \rightarrow \langle [Emb(\sigma)]_k \rangle_b$$
- Assume that the first L layers are simulated by p , we extend p to simulate layer $L + 1$
- If layer $L + 1$ is a FF layer, it's translated into B-RASP using position-wise operations
- If layer $L + 1$ is a SA layer, it's translated into B-RASP using attention operations

Translating transformers to B-RASP

Lemma

For any masked hard-attention transformer T , there is a B-RASP program p_T that simulates T .

Proof sketch

- The embedding function is simulated by
$$E_{k,b}(i) := \bigwedge_{\sigma \in \Sigma} Q_{\sigma}(i) \rightarrow \langle [Emb(\sigma)]_k \rangle_b$$
- Assume that the first L layers are simulated by p , we extend p to simulate layer $L + 1$
- If layer $L + 1$ is a FF layer, it's translated into B-RASP using position-wise operations
- If layer $L + 1$ is a SA layer, it's translated into B-RASP using attention operations
- Lastly, we add into the program an output vector Y to simulate T 's output layer

Translating B-RASP to transformers

Definition

Let p be a B-RASP program with vectors P_1, \dots, P_d . The transformer T (with $d' \geq d$ dimensions) simulates p iff for every word w of length n , for all $i \in [n], k \in [d]$ it holds

$$[T(w)]_{i,k} = \begin{cases} 1 & P_k(i) = 1 \\ 0 & \text{otherwise} \end{cases}$$

Translating B-RASP to transformers

Definition

Let p be a B-RASP program with vectors P_1, \dots, P_d . The transformer T (with $d' \geq d$ dimensions) simulates p iff for every word w of length n , for all $i \in [n], k \in [d]$ it holds

$$[T(w)]_{i,k} = \begin{cases} 1 & P_k(i) = 1 \\ 0 & \text{otherwise} \end{cases}$$

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Translating B-RASP to transformers

Definition

Let p be a B-RASP program with vectors P_1, \dots, P_d . The transformer T (with $d' \geq d$ dimensions) simulates p iff for every word w of length n , for all $i \in [n], k \in [d]$ it holds

$$[T(w)]_{i,k} = \begin{cases} 1 & P_k(i) = 1 \\ 0 & \text{otherwise} \end{cases}$$

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch

The proof is by induction of the number of vectors of p . Assume that the vectors P_1, \dots, P_k can be simulated by T_p . We show that P_{k+1} can be simulated as well.

Translating B-RASP to transformers: base cases

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

- Case P_{k+1} is an initial boolean vector for $\sigma \in \Sigma$.
we set $Emb(\sigma) = (0, \dots, 0, 1, 0, \dots, 0)$ where the 1 is at position $k + 1$

Translating B-RASP to transformers: base cases

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

- Case P_{k+1} is an initial boolean vector for $\sigma \in \Sigma$.
we set $Emb(\sigma) = (0, \dots, 0, 1, 0, \dots, 0)$ where the 1 is at position $k + 1$
- Case $P_{k+1}(i)$ is $BoolCombination(\{P_1(i), \dots, P_k(i)\})$.
It can be put in DNF and computed by the two layer FFN.

Translating B-RASP to transformers: base cases

Lemma


For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

- Case P_{k+1} is an initial boolean vector for $\sigma \in \Sigma$.
we set $Emb(\sigma) = (0, \dots, 0, 1, 0, \dots, 0)$ where the 1 is at position $k + 1$
- Case $P_{k+1}(i)$ is BoolCombination($\{P_1(i), \dots, P_k(i)\}$).

Example:

$$(a_1 \wedge \neg a_2 \wedge \neg a_3) \vee (b_1 \wedge b_2)$$



Translating B-RASP to transformers: base cases

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .


Proof sketch (cont.)

- Case P_{k+1} is an initial boolean vector for $\sigma \in \Sigma$.
we set $Emb(\sigma) = (0, \dots, 0, 1, 0, \dots, 0)$ where the 1 is at position $k + 1$
- Case $P_{k+1}(i)$ is $BoolCombination(\{P_1(i), \dots, P_k(i)\})$.

Example:

$$(a_1 \wedge \neg a_2 \wedge \neg a_3) \vee (b_1 \wedge b_2)$$

The first layer computes conjunctions



Translating B-RASP to transformers: base cases

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .


Proof sketch (cont.)

- Case P_{k+1} is an initial boolean vector for $\sigma \in \Sigma$.
we set $Emb(\sigma) = (0, \dots, 0, 1, 0, \dots, 0)$ where the 1 is at position $k + 1$
- Case $P_{k+1}(i)$ is $BoolCombination(\{P_1(i), \dots, P_k(i)\})$.

Example:

$$(a_1 \wedge \neg a_2 \wedge \neg a_3) \vee (b_1 \wedge b_2)$$

The second layer computes disjunctions



Translating B-RASP to transformers: attention operations

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : D(i).$$

Translating B-RASP to transformers: attention operations

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : D(i).$$

- Input vector: $(P_1(i), \dots, P_k(i), 0^{d-k}, 0, 0, \dots)$

Translating B-RASP to transformers: attention operations

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : D(i).$$

- Input vector: $(P_1(i), \dots, P_k(i), 0^{d-k}, 0, 0, \dots)$
- First layer: The SA sub-layer does nothing (identity function). The FFN adds (at free positions) information in order to compute $S(i,j)$

Translating B-RASP to transformers: attention operations

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] V(i,j) : D(i).$$

- Input vector: $(P_1(i), \dots, P_k(i), 0^{d-k}, 0, 0, \dots)$
- First layer: The SA sub-layer does nothing (identity function). The FFN adds (at free positions) information in order to compute $S(i,j)$
- Second Layer: The SA uses mask M and computes either $V(i,j')$ or $D(i)$. The FFN copies the desired result to position $k+1$.

Translating B-RASP to transformers: attention operations

Lemma

For any B-RASP program p , there is a transformer T_p that simulates p .

Proof sketch (cont.)

$$P_{k+1}(i) := \min/\max_j [M(i,j) \wedge S(i,j)] \quad V(i,j) : D(i).$$

- Input vector: $(P_1(i), \dots, P_k(i), 0^{d-k}, 0, 0, \dots)$
- First layer: The SA sub-layer does nothing (identity function). The FFN adds (at free positions) information in order to compute $S(i,j)$
- Second Layer: The SA uses mask M and computes either $V(i,j')$ or $D(i)$. The FFN copies the desired result to position $k+1$.
- Output vector: either $(P_1(i), \dots, P_k(i), V(i,j'), 0^{d-k-1}, \dots)$ or $(P_1(i), \dots, P_k(i), D(i), 0^{d-k-1}, \dots)$

B-RASP and Star Free Languages

B-RASP \rightarrow LTL

B-RASP \leftarrow Automata Cascades

LTL: definition

Syntax

$$\begin{aligned} \phi &:= p \mid \neg \phi \mid \phi \vee \phi \\ &\mid X\phi \mid \phi U \phi \\ &\mid Y\phi \mid \phi S \phi \end{aligned}$$

Boolean operators

Future operators

Past operators

LTL: definition

Syntax

$$\phi := p \mid \neg\phi \mid \phi \vee \phi$$

Boolean operators

$$\mid X\phi \mid \phi U \phi$$

Future operators

$$\mid Y\phi \mid \phi S \phi$$

Past operators

Semantics

Let \mathcal{AP} be a set of atomic propositions, and let $\sigma \in (2^{\mathcal{AP}})^+$ be a finite trace. Satisfaction of ϕ by σ at time $0 \leq i < |\sigma|$ is defined as follows.

LTL: definition

Syntax

$$\phi := p \mid \neg\phi \mid \phi \vee \phi$$

Boolean operators

$$\mid X\phi \mid \phi U \phi$$

Future operators

$$\mid Y\phi \mid \phi S \phi$$

Past operators

Semantics

Let \mathcal{AP} be a set of atomic propositions, and let $\sigma \in (2^{\mathcal{AP}})^+$ be a finite trace. Satisfaction of ϕ by σ at time $0 \leq i < |\sigma|$ is defined as follows.

Boolean operators

- $\sigma, i \models p$ iff $p \in \sigma_i$;
- $\sigma, i \models \neg\phi$ iff $\sigma, i \not\models \phi$;
- $\sigma, i \models \phi_1 \vee \phi_2$ iff $\sigma, i \models \phi_1$ or $\sigma, i \models \phi_2$;

LTL: definition

Syntax

$$\phi := p \mid \neg\phi \mid \phi \vee \phi$$

Boolean operators

$$\mid X\phi \mid \phi U \phi$$

Future operators

$$\mid Y\phi \mid \phi S \phi$$

Past operators

Semantics

Let \mathcal{AP} be a set of atomic propositions, and let $\sigma \in (2^{\mathcal{AP}})^+$ be a finite trace. Satisfaction of ϕ by σ at time $0 \leq i < |\sigma|$ is defined as follows.

Future operators

- $\sigma, i \models X\phi$ iff $i + 1 < |\sigma|$ and $\sigma, i + 1 \models \phi$;
- $\sigma, i \models \phi_1 U \phi_2$ iff there exists $i \leq j < |\sigma|$ such that $\sigma, j \models \phi_2$, and $\sigma, k \models \phi_1$ for all k , with $i \leq k < j$;

LTL: definition

Syntax

$$\phi := p \mid \neg\phi \mid \phi \vee \phi$$

Boolean operators

$$\mid X\phi \mid \phi U \phi$$

Future operators

$$\mid Y\phi \mid \phi S \phi$$

Past operators

Semantics

Let \mathcal{AP} be a set of atomic propositions, and let $\sigma \in (2^{\mathcal{AP}})^+$ be a finite trace. Satisfaction of ϕ by σ at time $0 \leq i < |\sigma|$ is defined as follows.

Past operators

- $\sigma, i \models Y\phi$ iff $i > 0$ and $\sigma, i - 1 \models \phi$;
- $\sigma, i \models \phi_1 S \phi_2$ iff there exists $j \leq i$ such that $\sigma, j \models \phi_2$, and $\sigma, k \models \phi_1$ for all k , with $j < k \leq i$;

LTL: shortcuts

We define the following shortcuts.

<i>(True)</i>	$True := p \vee \neg p$
<i>(False)</i>	$False := \neg True$
<i>(Weak Tomorrow)</i>	$\tilde{X}\phi := \neg X\neg\phi$
<i>(Eventually)</i>	$F\phi := True \ U \ \phi$
<i>(Globally)</i>	$G\phi := \neg F\neg\phi$
<i>(Weak Yesterday)</i>	$\tilde{Y}\phi := \neg Y\neg\phi$
<i>(Once)</i>	$O\phi := True \ S \ \phi$
<i>(Historically)</i>	$H\phi := \neg O\neg\phi$

Translating B-RASP to LTL

Lemma

For any Boolean vector P_k of a B-RASP program in normal form, there is a LTL formula ϕ_k such that for any input w of length n and all $i \in [n]$, it holds $P_k(i) = 1$ iff $w, i \models \phi_k$.

Translating B-RASP to LTL

Lemma

For any Boolean vector P_k of a B-RASP program in normal form, there is a LTL formula ϕ_k such that for any input w of length n and all $i \in [n]$, it holds $P_k(i) = 1$ iff $w, i \models \phi_k$.

Proof

Initial vectors and new vectors defined using point-wise operations are easily translated.

In the following, we show how to perform translation of vectors defined with attention operations starting with *max*. The cases with *min* are analogous.

Translating B-RASP to LTL: future masking

Lemma

For any Boolean vector P_k of a B-RASP program in normal form, there is a LTL formula ϕ_k such that for any input w of length n and all $i \in [n]$, it holds $P_k(i) = 1$ iff $w, i \models \phi_k$.

Proof (cont.)

$$P_k(i) := \max_j [j < i, S(j)] V(j) : D(i)$$

is translated to

$$\phi_k := Y(\neg\phi_S S(\phi_S \wedge \phi_V)) \vee ((\tilde{Y}H \neg\phi_S) \wedge \phi_D)$$

Formulas ϕ_S , ϕ_V and ϕ_D exist by induction hypothesis.

Observation

Future masking is translated using only past operators

Future masking

For any program in normal form, there is a sequence of length n and all $i \in [n]$, it

$$P_k(i) := \max_j [j < i, S(j)] V(j) : D(i)$$

is translated to

$$\phi_k := Y(\neg\phi_S \ S (\phi_S \wedge \phi_V)) \vee ((\tilde{Y}H \neg\phi_S) \wedge \phi_D)$$

Formulas ϕ_S , ϕ_V and ϕ_D exist by induction hypothesis.

Translating B-RASP to LTL: past masking

Lemma

For any Boolean vector P_k of a B-RASP program in normal form, there is a LTL formula ϕ_k such that for any input w of length n and all $i \in [n]$, it holds $P_k(i) = 1$ iff $w, i \models \phi_k$.

Proof (cont.)

$$P_k(i) := \max_j [j > i, S(j)] V(j) : D(i)$$

is translated to

$$\phi_k := XF(\phi_S \wedge \phi_V \wedge (\tilde{X}G \neg \phi_S)) \vee ((\tilde{X}G \neg \phi_S) \wedge \phi_D)$$

Formulas ϕ_S , ϕ_V and ϕ_D exist by induction hypothesis.

Observation

Past masking is translated using only future operators

masking

*program in normal form, there is
v of length n and all $i \in [n]$, it*

$$P_k(i) := \max_j [j > i, S(j)] V(j) : D(i)$$

is translated to

$$\phi_k := XF(\phi_S \wedge \phi_V \wedge (\tilde{X}G \neg \phi_S)) \vee ((\tilde{X}G \neg \phi_S) \wedge \phi_D)$$

Formulas ϕ_S , ϕ_V and ϕ_D exist by induction hypothesis.

Translating B-RASP to LTL: no masking

Lemma

For any Boolean vector P_k of a B-RASP program in normal form, there is a LTL formula ϕ_k such that for any input w of length n and all $i \in [n]$, it holds $P_k(i) = 1$ iff $w, i \models \phi_k$.

Proof (cont.)

$$P_k(i) := \max_j [1, S(j)] V(j) : D(i)$$

is translated to

$$\text{exists}(\phi) := (F\phi) \vee (O\phi)$$


$$\text{rightmost}(\phi, \psi) := \phi \wedge \psi \wedge (\tilde{X}G \neg\phi)$$

$$\phi_k := \text{exists}(\text{rightmost}(\phi_S, \phi_V)) \vee (\neg \text{exists}(\phi_S) \wedge \phi_D)$$

Formulas ϕ_S , ϕ_V and ϕ_D exist by induction hypothesis. \square

Automata Cascades and Krohn-Rhodes theory

Cascade product of automata is the generalization of the direct product




$$\mathcal{A} \circ \mathcal{B} = (\Sigma, Q_{\mathcal{A}} \times Q_{\mathcal{B}}, \delta) \text{ where}$$

$$\delta((q_{\mathcal{A}}, q_{\mathcal{B}}), \sigma) = (\delta_{\mathcal{A}}(q_{\mathcal{A}}, \sigma), \delta_{\mathcal{B}}(q_{\mathcal{B}}, (q_{\mathcal{A}}, \sigma)))$$

The second automaton reads the current state of the first one

Automata Cascades and Krohn-Rhodes theory

Cascade product of automata is the generalization of the direct product




$$\mathcal{A} \circ \mathcal{B} = (\Sigma, Q_{\mathcal{A}} \times Q_{\mathcal{B}}, \delta) \text{ where}$$
$$\delta((q_{\mathcal{A}}, q_{\mathcal{B}}), \sigma) = (\delta_{\mathcal{A}}(q_{\mathcal{A}}, \sigma), \delta_{\mathcal{B}}(q_{\mathcal{B}}, (q_{\mathcal{A}}, \sigma)))$$

The second automaton reads the current state of the first one

Lemma

Every counter-free automaton \mathcal{A} can be decomposed into a cascade $\mathcal{B}_1 \circ \dots \circ \mathcal{B}_n$ where each component \mathcal{B}_i is a two-state reset automaton


Two-states reset automaton



A transition labelled with σ can be of two types:

- Identity transition: $\forall q \ \delta(q, \sigma) = q$
- Reset transition on state q : $\forall q' \ \delta(q', \sigma) = q$

Two-states reset automaton



A transition labelled with σ can be of two types:

- Identity transition: $\forall q \ \delta(q, \sigma) = q$
- Reset transition on state q : $\forall q' \ \delta(q', \sigma) = q$

The alphabet Σ can be partitioned into sets of symbols I, R_{q_0}, R_{q_1} which respectively induce identity transitions, reset transitions entering state q_0 and reset transitions entering state q_1 . We denote with R the set of symbols which induce a reset transition on some state q .

B-RASP state simulation

Let \mathcal{A} be an automaton and let $w = w_1 \dots w_n$ be an input word. Let τ be the run s_0, \dots, s_n induced by w on \mathcal{A} starting from a certain state s .

B-RASP state simulation

Let \mathcal{A} be an automaton and let $w = w_1 \dots w_n$ be an input word. Let τ be the run s_0, \dots, s_n induced by w on \mathcal{A} starting from a certain state s . By definition of the run we have $s_{i+1} = \delta(s_i, w_i)$. s_i is the state reached before reading the i -th symbol.

B-RASP state simulation

Let \mathcal{A} be an automaton and let $w = w_1 \dots w_n$ be an input word. Let τ be the run s_0, \dots, s_n induced by w on \mathcal{A} starting from a certain state s . By definition of the run we have $s_{i+1} = \delta(s_i, w_i)$. s_i is the state reached before reading the i -th symbol.

Definition

A B-RASP program p is said to *simulate* \mathcal{A} starting from state s iff

- for each state q there is a boolean vector P_q ;
- for every input word w , $P_q(i) = 1$ iff $s_i = q$.

B-RASP from state simulation to language recognition

Given a program p which simulates the automaton \mathcal{A} , for every state r we define the predicate that decides whether \mathcal{A} started in state s ends up in state r after reading the symbol at position i :

$$A_r(i) := \bigvee_{\substack{q \in Q \\ \sigma \in \Sigma \\ \delta(q, \sigma) = r}} P_q(i) \wedge Q_\sigma(i)$$

B-RASP from state simulation to language recognition

Given a program p which simulates the automaton \mathcal{A} , for every state r we define the predicate that decides whether \mathcal{A} started in state s ends up in state r after reading the symbol at position i :

$$A_r(i) := \bigvee_{\substack{q \in Q \\ \sigma \in \Sigma \\ \delta(q, \sigma) = r}} P_q(i) \wedge Q_\sigma(i)$$

Given a set of final states F we can define the output vector of the program as

$$Y := \bigvee_{r \in F} A_r$$

and so $w \in \mathcal{L}(p)$ iff $w \in \mathcal{L}(\mathcal{A})$

Translating cascades to B-RASP

Lemma

Every cascade $\mathcal{C} = \mathcal{B}_1 \circ \dots \circ \mathcal{B}_k$ of two-states reset automata can be simulated by a B-RASP program $p_{\mathcal{C}}$ from state (s_1, \dots, s_k) .

Translating cascades to B-RASP

Lemma

Every cascade $\mathcal{C} = \mathcal{B}_1 \circ \dots \circ \mathcal{B}_k$ of two-states reset automata can be simulated by a B-RASP program $p_{\mathcal{C}}$ from state (s_1, \dots, s_k) .

Proof

We will proceed by induction on k the height of the cascade. The base case is a two-states reset automaton. For the inductive case, assume that the automaton $\mathcal{A} = \mathcal{B}_1 \circ \dots \circ \mathcal{B}_k$ can be simulated by the program $p_{\mathcal{A}}$ and show that $\mathcal{C} = \mathcal{A} \circ \mathcal{B}_{k+1}$ can be simulated by the program $p_{\mathcal{C}}$.

Translating cascades to B-RASP: base case

Lemma

Every cascade $\mathcal{C} = \mathcal{B}_1 \circ \dots \circ \mathcal{B}_k$ of two-states reset automata can be simulated by a B-RASP program $p_{\mathcal{C}}$ from state (s_1, \dots, s_k) .

Proof (cont.)

Let $\mathcal{A} = (\Sigma, Q, \delta)$ be a reset automaton. The following program $p_{\mathcal{A}}$, simulates \mathcal{A} starting from state $s \in Q$. Q_{σ} denotes the input vector of σ .

$$P_q(i) := \max_j \left[j < i, \bigvee_{\sigma \in R} Q_{\sigma}(j) \right] \bigvee_{\sigma \in R_q} Q_{\sigma}(j) : 0 \quad \text{if } q \neq s$$

$$P_q(i) := \max_j \left[j < i, \bigvee_{\sigma \in R} Q_{\sigma}(j) \right] \bigvee_{\sigma \in R_q} Q_{\sigma}(j) : 1 \quad \text{if } q = s$$

Translating cascades to B-RASP: inductive case

Lemma

Every cascade $\mathcal{C} = \mathcal{B}_1 \circ \dots \circ \mathcal{B}_k$ of two-states reset automata can be simulated by a B-RASP program $p_{\mathcal{C}}$ from state (s_1, \dots, s_k) .

Proof (cont.)

The automaton $\mathcal{A} = (\Sigma, Q_{\mathcal{A}}, \delta_{\mathcal{A}})$ is simulated by $p_{\mathcal{A}}$. $P_{q_{\mathcal{A}}}$ is the predicate for state $q_{\mathcal{A}}$.

The reset automaton $\mathcal{B} = (\Sigma \times Q, Q_{\mathcal{B}}, \delta_{\mathcal{B}})$ is simulated by $p_{\mathcal{B}}$. $P_{q_{\mathcal{B}}}$ is the predicate for state $q_{\mathcal{B}}$.

Define $Q'_{(q_{\mathcal{A}}, \sigma)} = P_{q_{\mathcal{A}}} \wedge Q_{\sigma}$ and $P'_{q_{\mathcal{B}}}$ as a clone of $P_{q_{\mathcal{B}}}$ where $Q_{(q_{\mathcal{A}}, \sigma)}$ is replaced by $Q'_{(q_{\mathcal{A}}, \sigma)}$.

For a state $(s_{\mathcal{A}}, s_{\mathcal{B}}) \in Q_{\mathcal{A}} \times Q_{\mathcal{B}}$ of $\mathcal{C} = \mathcal{A} \circ \mathcal{B}$ we define the predicate $P_{(s_{\mathcal{A}}, s_{\mathcal{B}})}$ as $P_{(s_{\mathcal{A}}, s_{\mathcal{B}})} := P_{s_{\mathcal{A}}} \wedge P'_{s_{\mathcal{B}}}$ □

Conclusions

Further results

Well known results of LTL have deep consequences on transformers

Lemma

Pure Past LTL (pLTL) is equivalent to LTL

Lemma

LTL with non-strict operators recognize exactly the stutter-invariant star-free languages

Lemma

There exists a language L_k such that no LTL formula with temporal depth $2k$ recognizes L_k , but there exists an LTL formula with temporal depth $2k + 1$ which recognizes L_k

Further results

Well known results of LTL have deep consequences on transformers

Lemma

Transformers with only future-masked rightmost-hard attention recognize exactly the star-free languages.

Lemma

LTL with non-strict operators recognize exactly the stutter-invariant star-free languages

Lemma

There exists a language L_k such that no LTL formula with temporal depth $2k$ recognizes L_k , but there exists an LTL formula with temporal depth $2k + 1$ which recognizes L_k

Further results

Well known results of LTL have deep consequences on transformers

Lemma

Transformers with only future-masked rightmost-hard attention recognize exactly the star-free languages.

Lemma

Masked hard-attention transformers with only non-strict masking recognize exactly the stutter-invariant star-free languages.

Lemma

There exists a language L_k such that no LTL formula with temporal depth $2k$ recognizes L_k , but there exists an LTL formula with temporal depth $2k + 1$ which recognizes L_k

Further results

Well known results of LTL have deep consequences on transformers

Lemma

Transformers with only future-masked rightmost-hard attention recognize exactly the star-free languages.

Lemma

Masked hard-attention transformers with only non-strict masking recognize exactly the stutter-invariant star-free languages.

Lemma

There exists a language L_k such that no multi-head masked hard-attention transformer of depth k recognizes L_k , but there exists a transformer of depth $O(k)$ which recognizes L_k

Depth Hierarchy

This last result is in contrast to the universal approximation theorem for standard FFN

sequences on transformers

most-hard attention recognize

Masked hard-attention transformers with only non-strict masking recognize exactly the stutter-invariant star-free languages.

Lemma

There exists a language L_k such that no multi-head masked hard-attention transformer of depth k recognizes L_k , but there exists a transformer of depth $O(k)$ which recognizes L_k

Conclusions and final comments

● Recap

- ▶ B-RASP has been used as an intermediate language which facilitates translation to LTL
- ▶ Both SA and FF sub-layers of a transformer play a significant role. FF corresponds to boolean operators while SA corresponds to temporal operators

● Future Work

- ▶ Extend the study to soft-attention transformers
- ▶ Extend the study to encoder-decoder transformers

● My Questions

- ▶ How does training transformers compare to automata learning algorithms?
- ▶ LTL is used to derive results about transformers: is the converse possible?
- ▶ How do reset automata cascades connect to transformer expressiveness?

Bibliography

Main paper



Andy Yang, David Chiang, and Dana Angluin. *Masked Hard-Attention Transformers Recognize Exactly the Star-Free Languages*. 2024. URL:
<https://arxiv.org/abs/2310.13897>

LTL reference



Geatti Luca. *Temporal Logic Specifications: Expressiveness, Satisfiability And Realizability*. 2022

Automata Cascades reference



Borelli Roberto, Geatti Luca, Montali Marco, and Montanari Angelo. *On Cascades of Reset Automata*. 2025